ros2ai Next-Gen ROS 2 CLI

ros2ai is a next-generation ROS 2 command line interface extension with OpenAI, Anthropic and Ollama

by **fujitatomoya@github** / **tomoyafujita@linkedin**

Everything Here!!!

Issues/PRs Welcome



by ***fujitatomoya@github*** / ***tomoyafujita@linkedin***


0:00 / 0:24


0:00 / 0:33


0:00 / 0:24

Motivation

- (Originally just for fun 😂)
- Quickly get the answers against questions.
- Multiple Language Support.
- Support beginners and students. (your best trainer)
- Easy to use for everyone.
- Dedicated / AI based Support (your own concierge)
- Bridge / Proxy to LLM
- Agnostic from LLM service (OpenAI, Anthropic, Ollama or any OpenAI compatible endpoint)

Background

I want the answers w/o searching...

browsing, clicking and typing many times to get to the information does not work for me. I need ***what***, not ***where***. all these are small things, but can be easily compiled up to mountain, especially beginners.

I really do not care how to use ros2cli...

many sub-commands, options, and arguments. besides, those could be deprecated or removed time to time... this does not work for me. actually i just want to ask "what parameters are available?", "check the details for the topic /chatter", but before that i need to know how to do that... this is already barrier for beginners.

it would be better for maintenance if someone abstracts this to the user, so that maintainers can just go ahead to remove, change options w/o deprecation or certain soak time...

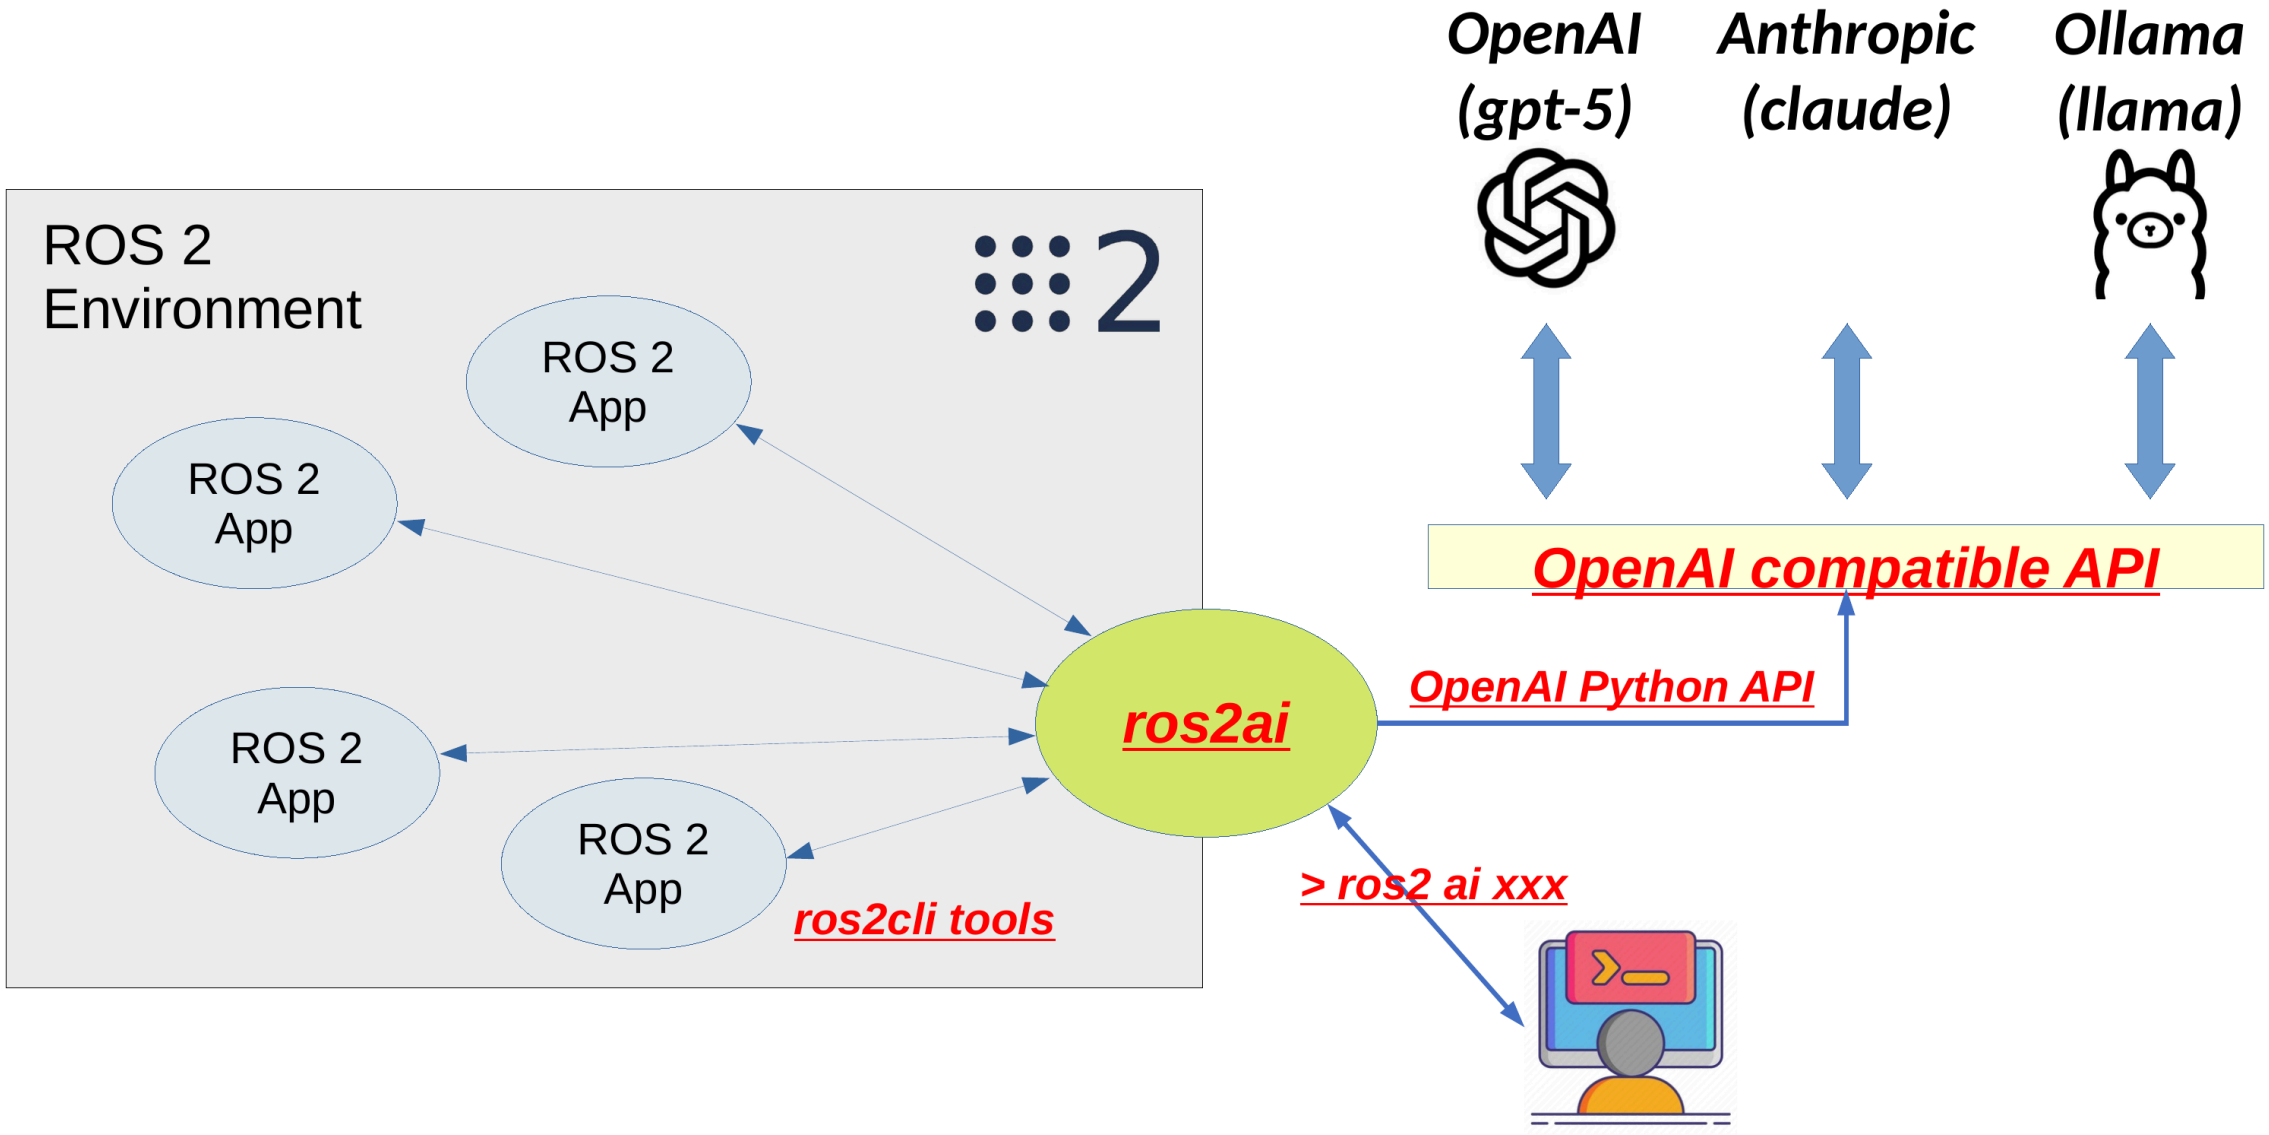
Multi-Language Documentation

OSS projects tend to add multi-language support documentation under mainline doc. this is really likely, and they do. but the question is ***Does it really work?***... IMO, it does (will) not. eventually those docs will be outdated, not maintained, not scalable, not precise and once that happens these docs would be garbage.

i believe that is exactly where AI comes to play. instead of paying the resources to non-scalable things, focus on the mainline with precise information which is available for any languages.

Design

- ***SIMPLE***. one of the original motivation, it has to be simple as much as possible. That is said that there is no sub-command at all, just `ros2 ai <whatever your request>`.



by fujitatomoya@github / tomoyafujita@linkedin

- `ros2 ai "<request>"` is the only entry point, the request can be any question related to ROS 2, or the operation against the ROS 2 system. LLM system role attribution is set to default, but can be reconfigurable.
- AI decides by itself whether to answer the question, or to inspect and operate the ROS 2 system via Function Calling with ros2cli tools. The commands that change the system require the user confirmation.
- `ros2 ai --status` is to check if `ros2ai` is configured properly.
- This is a single-shot completion, no session is supported currently.

Tool Calling Policy

- Every tool maps to ros2cli, executed without shell and printed before execution.
- **read-only** (`node`, `topic list/info/echo`, `param get`...) run automatically.
- **mutating** (`topic pub`, `service call`, `param set`...) require confirmation, `--yes` skips it, `--dry-run` only prints the candidates.
- **long-running** (`run`, `launch`, `bag record`...) require confirmation and run in foreground until `Ctrl-C`.
- **composite** (`get_node_status`, `get_topic_status`) run several commands in one call and return a YAML summary with warnings.

Documentation Grounding

- The answers about ROS 2 concepts, tutorials, tools and APIs are grounded on the official ROS 2 documentation of the running distribution (`ROS_DISTRO`), and cite the page.
- `search_ros2_docs` and `read_ros2_docs` look up the documentation sources cached locally (`~/.cache/ros2ai/docs/<distro>`), downloaded once per distribution and refreshed weekly, offline afterwards.
- Multi-language support stays on the mainline documentation, the model translates on demand.

Future Plans

Model Comparison

- Experimental performance and accuracy comparison over different LLM models.

API Parameter Adjustment

- for more stable AI behavior.
- performance and accuracy comparison.
- latency for user experience.

Context-Awareness

currently, `ros2ai` only supports single-shot completion API, that means we can not rely on the previous questions or answers to make further requests. e.g) "Is /chatter topic available?", "Subscribe it!". This ***it*** should be recognized in the same context or session. If this command is issued, prompt should be initiated to keep the session alive until the entire session is closed by the user.

Audio Interface Support

It would be really helpful to iterate via audio system instead of typing anymore. So the user can just talk to the `ros2ai` agent and then `ros2ai` answers any questions and executes the corresponding operations as user wants to do.

Fine-tuning

- ROS 2 general fine-tuning
  - this requires scaled training dataset for ROS 2, agnostic from user environment. could be distro specific. so that AI can response more precisely based on questions and requests.
- User specified environment tuning
  - this dataset should be uploaded to help more user specific problems and questions. if this is doable, `ros2ai` agent can the 1st customer support for anyone?